

**Prerequisites: You must have external network connectivity.**

## DHCP ( neutron)

- When we create the subnet the DHCP get created automatically with the first IP of the subnet (192.168.122.100 here)

\$ neutron net-list // to check the network list (public/external network)

```
+-----+-----+-----+-----+
| id          | name      | tenant_id | subnets      |
+-----+-----+-----+-----+
| 6d5e541e-b439-4e7b-b7ee-451db6805c12 | public_network | 14fbf2321b1c45b682fd6835895c1e4d | 29f61c6a-3025-4b73-9238-ef81c41267e2 192.168.122.0/24 |
+-----+-----+-----+-----+

$ neutron port-create --fixed-ip subnet_id=29f61c6a-3025-4b73-9238-ef81c41267e2,ip_address=192.168.122.149 public_network
```

## Create subnet for OCP cluster and provide range for OCP hosts

# source overcloudrc

# neutron subnet-create --dns-nameserver 8.8.8.8,8.8.4.4 --allocation-pool  
start=192.168.122.100,end=192.168.122.150

*// We will change the DNS server later as we are creating it for our OCP cluster.*

# neutron subnet-list

neutron CLI is deprecated and will be removed in the future. Use openstack CLI instead.

```
+-----+-----+-----+-----+-----+
| id          | name      | tenant_id | cidr          | allocation_pools |
+-----+-----+-----+-----+-----+
| 29f61c6a-3025-4b73-9238-ef81c41267e2 | public_subnet | 14fbf2321b1c45b682fd6835895c1e4d | 192.168.122.0/24 | {"start": "192.168.122.100", "end": "192.168.122.150"} |
+-----+-----+-----+-----+-----+
```

**To get a fixed IP address, create ports with static IP for the OCP hosts.**

**# for bootstrap**

```
neutron port-create --fixed-ip  
subnet_id=29f61c6a-3025-4b73-9238-ef81c41267e2,ip_address=192.168.122.101  
public_network
```

**# for Master**

```
neutron port-create --fixed-ip  
subnet_id=29f61c6a-3025-4b73-9238-ef81c41267e2,ip_address=192.168.122.102  
public_network
```

**# for worker**

```
neutron port-create --fixed-ip  
subnet_id=29f61c6a-3025-4b73-9238-ef81c41267e2,ip_address=192.168.122.103  
public_network
```

**# for LB**

```
neutron port-create --fixed-ip  
subnet_id=29f61c6a-3025-4b73-9238-ef81c41267e2,ip_address=192.168.122.105  
public_network
```

## DNS:

Now we create the DNS server on 192.168.122.149 (again this is a fixed port we created it)

> Create RHEL DNS host ( using **dnsmasq**)

- Adding a records in /etc/addnhost

```
#####  
/etc/addnhost  
192.168.122.101 bootstrap.ocplab.example.com  
192.168.122.102 master.ocplab.example.com etcd-0.ocplab.example.com  
192.168.122.103 worker.ocplab.example.com  
192.168.122.149 lb.ocplab.example.com api.ocplab.example.com api-int.ocplab.example.com  
lb  
192.168.122.149 dns-server-ocp-lab.ocplab.example.com  
#####
```

- Config of dnsmasq:

```
~~~  
/etc/dnsmasq.d/abc  
local=/ocplab.example.com/  
address=/apps.ocplab.example.com/192.168.122.149  
srv-host=_etcd-server-ssl._tcp.ocplab.example.com,master.ocplab.example.com,2380,0,10  
no-hosts  
server=192.168.122.1  
addn-hosts=/etc/addnhosts  
~~~
```

**[IMP]\*\* Update DNS server in the subnet that we have created.**

```
# neutron subnet-update --dns-nameserver 192.168.122.149  
29f61c6a-3025-4b73-9238-ef81c41267e2
```

## LB (host haproxy)

- Sample haproxy config having 1 master configuration

```
#####
# Global settings
#-----
    log      127.0.0.1:514 local0 debug
    chroot    /var/lib/haproxy
    pidfile   /var/run/haproxy.pid
    maxconn   10000
    user      haproxy
    group     haproxy
    daemon

    # turn on stats unix socket
    stats socket /var/lib/haproxy/stats

#-----
# common defaults that all the 'listen' and 'backend' sections will
# use if not designated in their block
#-----
defaults
    mode                http
    log                  global
    option               httplog
    option               dontlognull
    option               redispatch
    retries              3
    timeout http-request 10s
    timeout queue        1m
    timeout connect      10s
    timeout client        300s
    timeout server       300s
```

```
timeout http-keep-alive 10s
timeout check          10s
maxconn                10000
```

```
#-----OCP API-----
```

```
frontend atomic-openshift-api
    bind *:6443
    default_backend atomic-openshift-api
    mode tcp
    option tcplog
```

```
backend atomic-openshift-api
    balance source
    mode tcp
    server bootstrap 10.74.252.114:6443 check
    server master 10.74.255.187:6443 check
```

```
#-----MACHINE CONFIGS-----
```

```
frontend machine-config-server
    bind *:22623
    default_backend machine-config-server
    mode tcp
    option tcplog
```

```
backend machine-config-server
    balance source
    mode tcp
    server bootstrap 10.74.252.114:22623 check
    server master 10.74.255.187:22623 check
```

```
#-----HTTP traffic-----
```

```
frontend http-forward
    bind *:80
    use_backend http-backend
    mode http
```

```
backend http-backend
    balance leastconn
    mode http
```

server            worker 10.74.255.187:80 check

#-----HTTPS traffic-----

frontend https-forward

bind \*:443

default\_backend https-backend

mode tcp

option tcplog

backend https-backend

balance leastconn

mode tcp

server            worker 10.74.255.187:443 check

#####

## BASTION:

- Download [openshift-install-linux-4.2.12.tar.gz](#) and [openshift-client-linux-4.2.12.tar.gz](#) and extract it.
- Create installer-config [Read](#)  
[https://docs.openshift.com/container-platform/4.2/installing/installing\\_bare\\_metal/installing\\_bare-metal.html#installation-initializing-manual\\_installing-bare-metal](https://docs.openshift.com/container-platform/4.2/installing/installing_bare_metal/installing_bare-metal.html#installation-initializing-manual_installing-bare-metal)
- Create manifests and ignition files [Read](#)  
[https://docs.openshift.com/container-platform/4.2/installing/installing\\_bare\\_metal/installing\\_bare-metal.html#installation-user-infra-generate-k8s-manifest-ignition\\_installing-bare-metal](https://docs.openshift.com/container-platform/4.2/installing/installing_bare_metal/installing_bare-metal.html#installation-user-infra-generate-k8s-manifest-ignition_installing-bare-metal)

## HTTP host (RHEL)

Download the [rhcos-4.2.0-x86\\_64-installer.iso](#) and [rhcos-4.2.0-x86\\_64-metal-bios.raw.gz](#) and place it in http server. Also, place the ignition configs in it.

## Boot the OCP Instances

### Create bootstrap host on OSP

- Download rhcos iso and make it available on openstack ( on over-cloud):

```
# openstack image create --container-format=bare --disk-format=iso  
--file=rhcos-4.2.0-x86_64-installer.iso --public rhcos
```

```
# openstack image list
```

- Create temporary instance using the newly created image:

```
# nova boot --flavor <flavor-name> --image rhcos <instance-name>
```

- Create a volume and add it to the instance:

```
# openstack volume create --size <size> --bootable <vol-name>
```

```
# openstack server add volume <instance-name> <vol-name>t --device /dev/vda
```

- On the OSP console, open up the temporary instance and check if the volume is attached to the instance. After that install the coreos on the attached volume:

```
# /usr/libexec/coreos-installer -d vda -i <path/to/ign/file> -b <path/to/raw.gz>
```

- After the installation is complete, delete the temporary instance and create the qcow2 image using the volume from that instance:

```
# openstack server delete <temporary-instance>
```

```
# openstack image create --volume <vol-name> <img-name>
```

- Create the node from newly created image:

```
# nova boot --flavor <flavor-name> --image <img-name> <name-of-node> --nic  
port-id=20d55d3-e018-4421-95dc-700ca7ac449 // which we created earlier
```

**NOTE:** Apply the aforementioned process for the master and worker nodes(hosts), just change the ignition configs as per the hosts

**IMP:** Why have we followed the above approach for creating coreOS bootable image?

- As OSP has constraint that it accepts only a single bootable disk image(iso or qcow2 or raw or etc..) to boot the instance.
- So we initially mount the ISO at the time of boot and it takes as Read Only partition, so for that we attached a separate volume for instance form cinder and mark that volume as "bootable".
- After that we will pass the aforementioned volume for coreos installation along with the other OCP coreos parameters the configs and data will be written in the Cinder Volume.
- Now using this volume we create a bootable qcow2 OSP compatible image and use this to boot the CoreOs instance for installation.

**After installation:**

Monitor the logs on master and workers

- Check the status of nodes

```
# ./oc get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master.ocplab.example.com	Ready	master	156m	v1.14.6+cebabbf4a
worker.ocplab.example.com	Ready	worker	156m	v1.14.6+cebabbf4a



## **KNOWN ISSUE:**

> After the installation the instances boot with the random names. Reset the names as we set in DNS and restart the kubelet service on all hosts.

```
# systemctl restart kubelet
```

> After that just approve the pending CSR's

```
# oc get csr -o name | xargs oc adm certificate approve
```